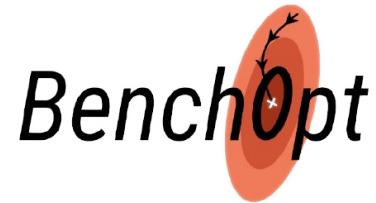


BenchOpt Tutorial



The ImageNet competition

- Annual competition since 2010



The ImageNet competition

- Annual competition since 2010
- Evaluate image classification methods with 14M labeled images among 1k categories



The ImageNet competition

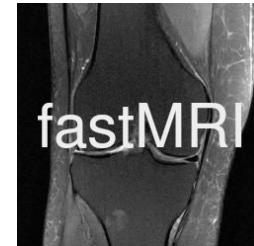
- Annual competition since 2010
- Evaluate image classification methods with 14M labeled images among 1k categories
- Boosted AI and Deep Learning research when Alex Krizhevsky won in 2012.



Many benchmarks

Many benchmarks followed ImageNet:

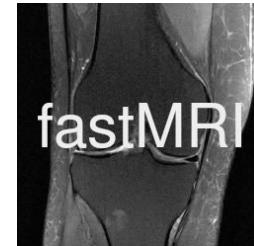
- Natural Language Processing:
GLUE, SuperGLUE
- Reinforcement Learning: Atari,
MuJoCo, OpenAI Gym
- Others: fastMRI, DAWNBench,
MLPerf, etc.



Many benchmarks

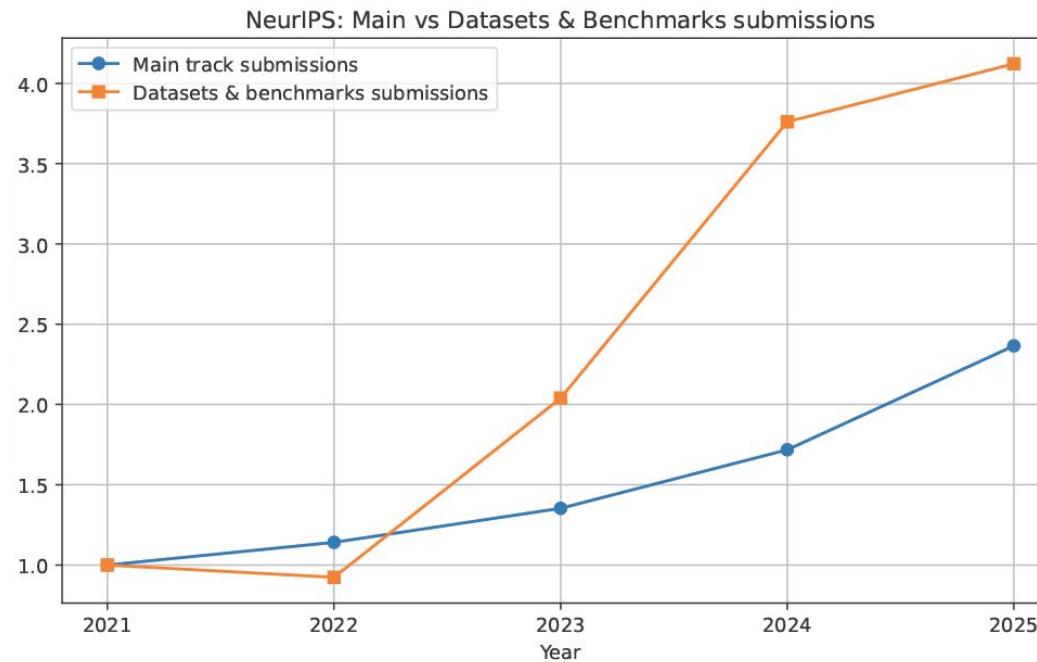
Many benchmarks followed ImageNet:

- Natural Language Processing:
GLUE, SuperGLUE
- Reinforcement Learning: Atari,
MuJoCo, OpenAI Gym
- Others: fastMRI, DAWNBench,
MLPerf, etc.



→ Benchmarks are now ubiquitous in AI research.

Too many benchmarks in AI?



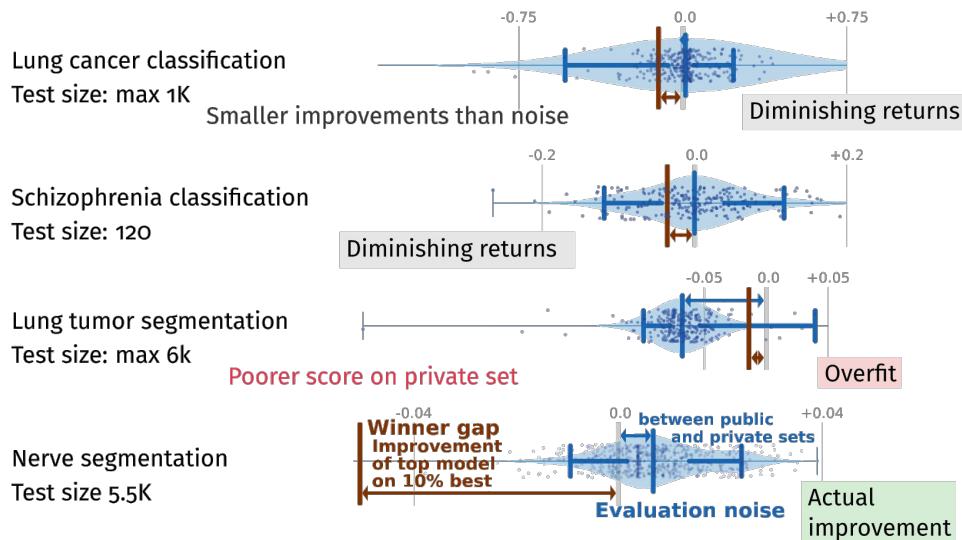
Benchmark goals in AI

	Short term progress	Long term progress
Task specific	Challenge/Competition → push limits quickly	SOTA tracking → measure progress
Generalizable	Research question → empirical study	Benchmarking framework → stable & extensible

- Most attention goes to the top-left quadrant for fast progress, but solid science requires the bottom-right.

Challenges in benchmarking AI

- Futile benchmarks



[Varoquaux and Cheplygina 2022]

Challenges in benchmarking AI

- Futile benchmarks
- Lack of proper baselines hinders scientific progress.

Do we really need Foundation Models for multi-step-ahead Epidemic Forecasting?

Position: Quo Vadis, Unsupervised Time Series Anomaly Detection?

M. Saquib Sarfraz^{1,2} Mei-Yen Chen¹ Lukas Layer¹ Kunyu Peng² Marios Koulakis²

Descending through a Crowded Valley —
Benchmarking Deep Learning Optimizers

Robin M. Schmidt^{*1} Frank Schneider^{*1} Philipp Hennig^{1,2}

PNAS

RESEARCH ARTICLE | COMPUTER SCIENCES

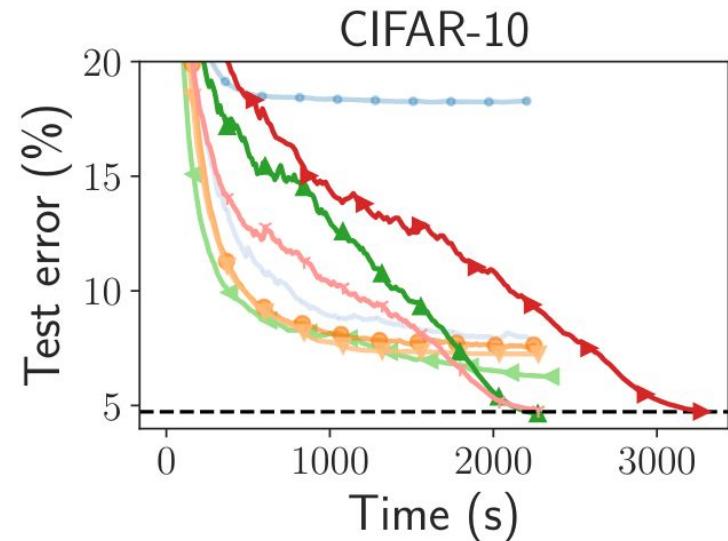
OPEN ACCESS

Implicit data crimes: Machine learning bias arising from misuse of public data

Efrat Shimron  a,¹ Jonathan I. Tamir  b,c,d, Ke Wang^a, and Michael Lustig^a

Challenges in benchmarking AI

- Futile benchmarks
- Lack of proper baselines hinders scientific progress.
- Reproducing benchmarks is hard.



[Moreau et al. 2022]

Challenges in benchmarking AI

- Futile benchmarks
- Lack of proper baselines hinders scientific progress.
- Reproducing benchmarks is hard.
- Statistical validity is often missing.

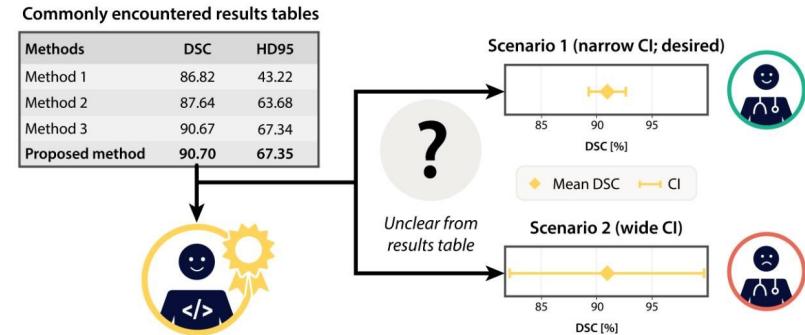
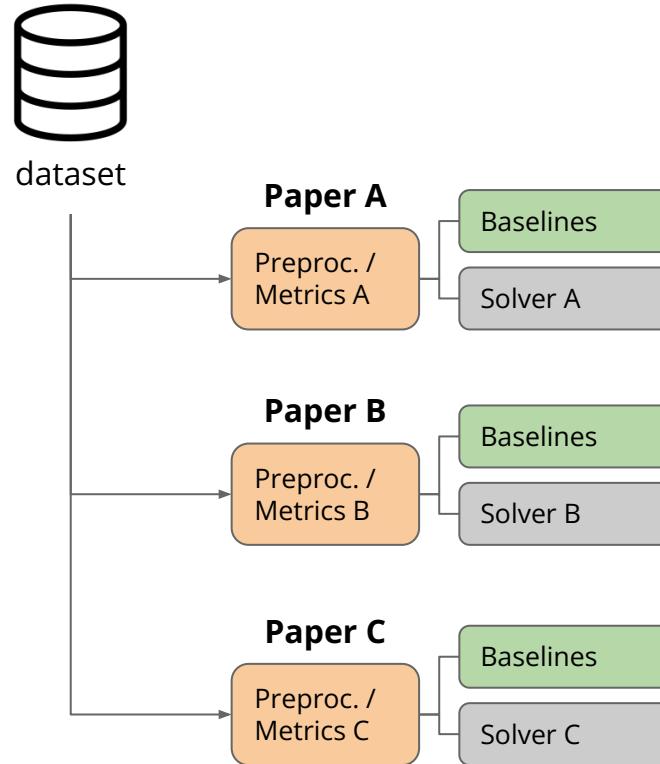


Fig. 1: Common practice in medical imaging algorithm performance reporting leaves many open questions.

[Christodoulou et al. 2024]

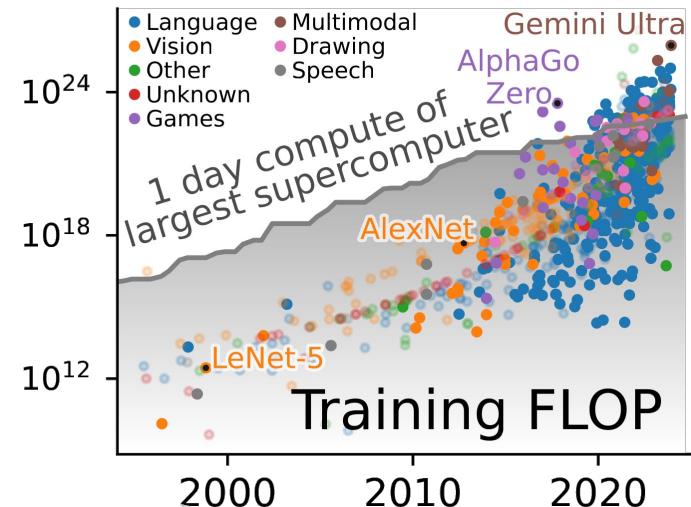
Challenges in benchmarking AI

- Futile benchmarks
- Lack of proper baselines hinders scientific progress.
- Reproducing benchmarks is hard.
- Statistical validity is often missing.
- Benchmarking cost is duplicated across groups.



Challenges in benchmarking AI

- Futile benchmarks
- Lack of proper baselines hinders scientific progress.
- Reproducing benchmarks is hard.
- Statistical validity is often missing.
- Benchmarking cost is duplicated across groups.



[Varoquaux et al. 2025]

Challenges in benchmarking AI

Technical challenges for repeatability:

- ▶ Multiple frameworks: PyTorch, Tensorflow, jax,...
- ▶ Hardware dependence: CPU, GPU, TPUs,...
- ▶ Large scale datasets: hard to share, pre-process,...

Challenges in benchmarking AI

Technical challenges for repeatability:

- ▶ Multiple frameworks: PyTorch, Tensorflow, jax,...
- ▶ Hardware dependence: CPU, GPU, TPUs,...
- ▶ Large scale datasets: hard to share, pre-process,...

Challenges for reusability:

- ▶ Usability: API, modularity, documentation, tests,...
- ▶ Maintenance: dependencies, code rot, long term support,...

Challenges in benchmarking AI

Technical challenges for repeatability:

- ▶ Multiple frameworks: PyTorch, Tensorflow, jax,...
- ▶ Hardware dependence: CPU, GPU, TPUs,...
- ▶ Large scale datasets: hard to share, pre-process,...

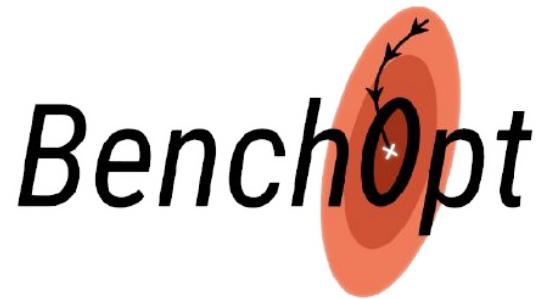
Challenges for reusability:

- ▶ Usability: API, modularity, documentation, tests,...
- ▶ Maintenance: dependencies, code rot, long term support,...

Statistical challenges for replicability:

- ▶ Stochastic algorithms: random initialization, data shuffling,...
- ▶ Data handling: splitting, data preprocessing, model selection,...

Running benchmarks with Benchopt



Benchopt provides a framework to organize and run benchmarks

- Moreau, Thomas, et al. (2022) "Benchopt: Reproducible, efficient and collaborative optimization benchmarks." In *NeurIPS*

Running benchmarks with Benchopt

Examples of existing benchmarks (<https://github.com/benchopt>):

- Image Classification (resnet)
- Logistic regression
- Lasso
- ICA
- Unsup. Domain Adaptation
- Bilevel Optimization
- Brain Computer Interface
- ...

Running benchmarks with Benchopt

Clone the repo https://github.com/Etyl/benchmark_classification

Create virtual environment and install dependencies:

With conda:

```
conda create -n benchopt-clf python=3.11
conda activate benchopt-clf
python -m pip install benchopt
benchopt install . --env-name benchopt-clf
```

With venv:

```
python -m venv venv
source venv/bin/activate
python -m pip install -r requirements.txt
```

Running benchmarks with Benchopt

Running the benchmark

```
benchopt run .
```

Running benchmarks with Benchopt

Running the benchmark

```
benchopt run .
```

Running the benchmark with a config

```
benchopt run . --config run_config.yml
```

```
# run_config.yml

dataset:
    simulated

solver:
    svm:
        kernel: ["linear"]
    torch-logreg:
        lr: [0.01]
        batch_size: [64, 256]

n_repetitions: 3
timeout: 30
seed: 0
n-jobs: 4
```

Running benchmarks with Benchopt

Running the benchmark

```
benchopt run .
```

Running the benchmark with a config

```
benchopt run . --config run_config.yml
```

```
# run_config.yml

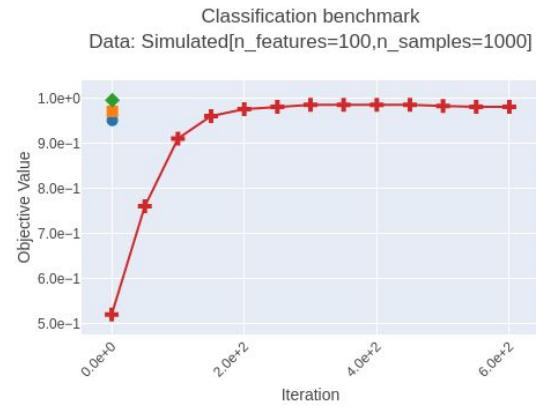
dataset:
    simulated

solver:
    svm:
        kernel: ["linear"]
    torch-logreg:
        lr: [0.01]
        batch_size: [64, 256]

n_repetitions: 3
timeout: 30
seed: 0
n-jobs: 4
```

- results are cached when running the same setup

Interactive results exploration

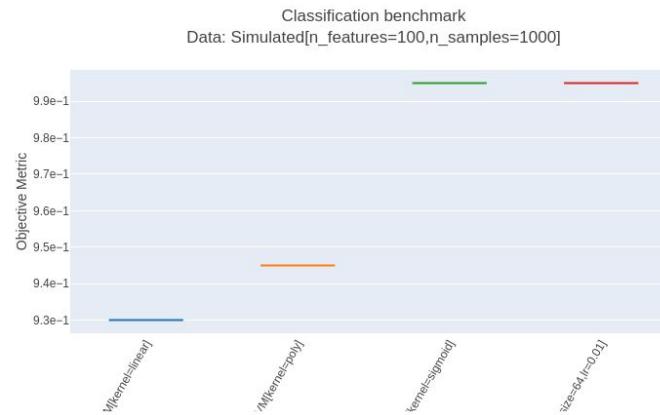


Solvers (Click on a solver to hide it or double click to hide all the others)

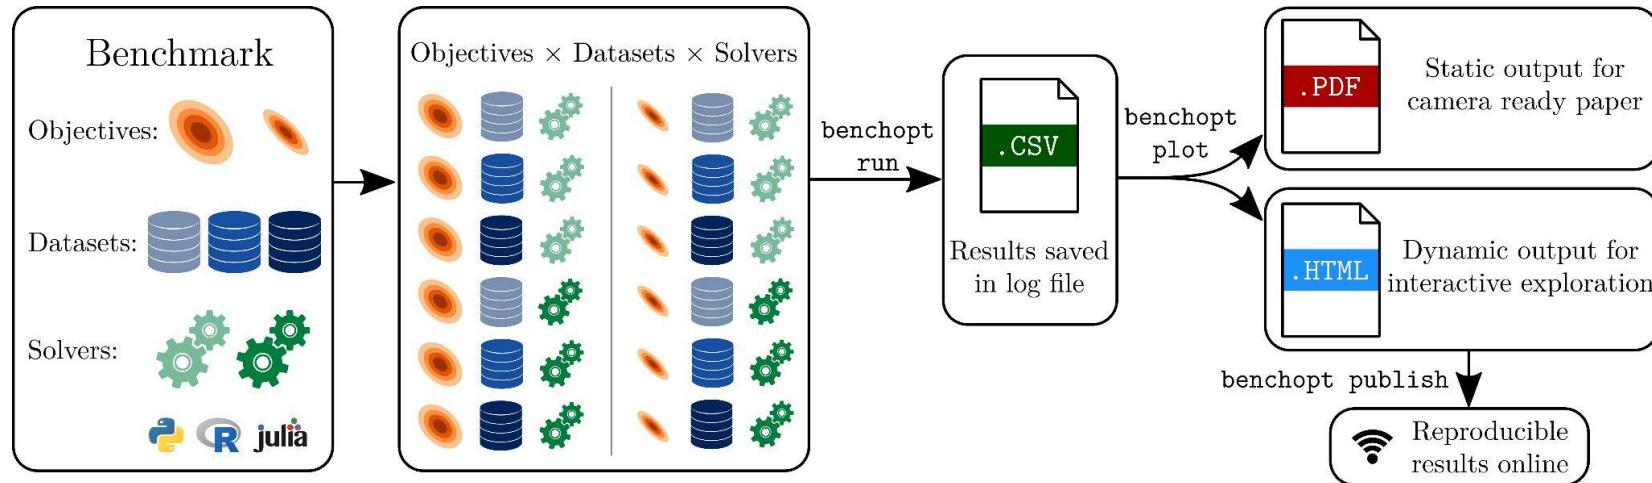
- SVM[kernel=linear]
- SVM[kernel=poly]
- SVM[kernel=sigmoid]
- torch-logreg[batch_size=64,lr=0.01]

solver	accuracy_test	accuracy_train	time (s)
SVM[kernel=linear]	0.9500	1	0.0371
SVM[kernel=poly]	0.9700	1	0.0493
SVM[kernel=sigmoid]	0.9950	0.9888	0.0291
torch-logreg[batch_size=64,lr=0.01]	0.9800	0.9925	0.6374

Float Precision: 4 Export LaTeX

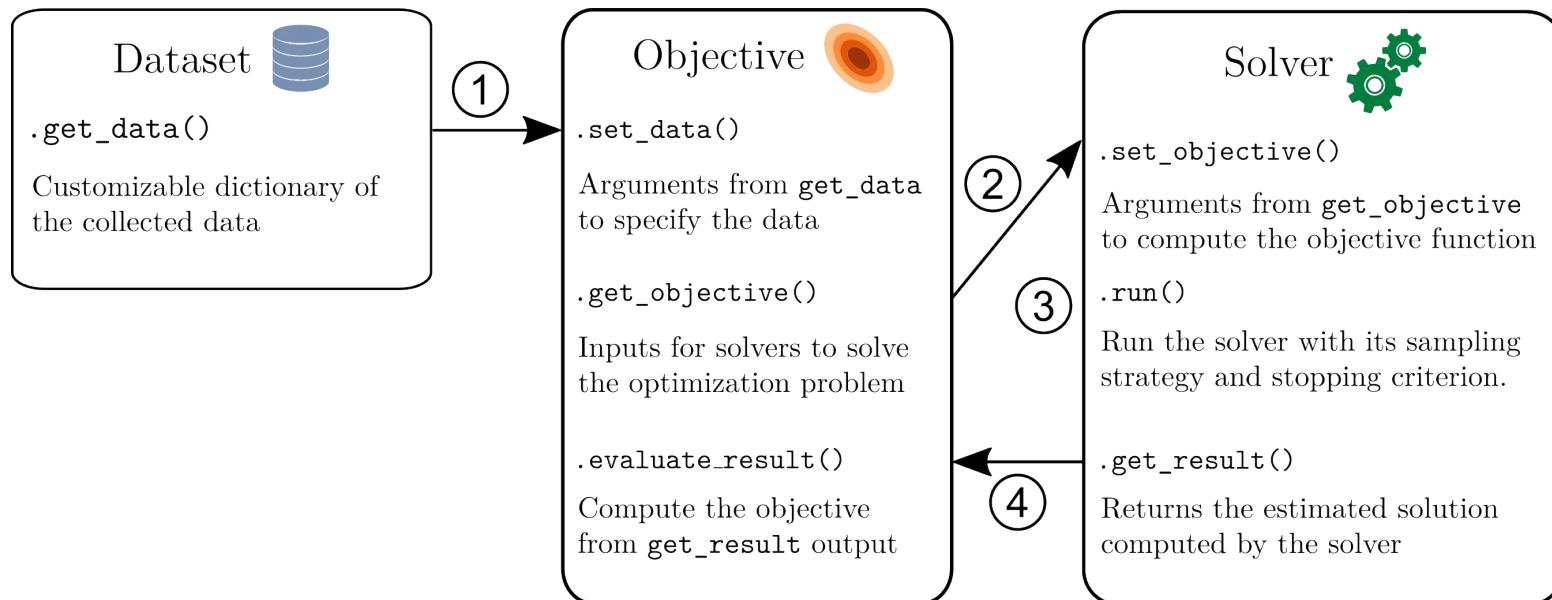


A modular framework to create benchmarks



A modular framework to create benchmarks

Dependency relation between Dataset - Objective - Solver



Structure of a benchmark

```
benchmark/
    └── objective.py
    └── datasets/
        ├── dataset1.py
        └── dataset2.py
    └── solvers/
        ├── solver1.py
        └── solver2.py
    └── plots/
        ├── custom_plot1.py
        └── custom_plot2.py
```

Adding a dataset

```
class Dataset(BaseDataset):
    name = "Simulated"
    parameters = {
        'n_samples, n_features': [
            (1000, 100),
            (5000, 10),
        ],
    }
    requirements = []

    def get_data(self):
        X, y = make_classification(
            n_samples=self.n_samples, n_features=self.n_features,
            n_informative=1, n_redundant=0, n_clusters_per_class=1,
            random_state=self.get_seed()
        )
        return dict(X=X, y=y)
```

Adding a dataset

Add a new dataset (for example the breast cancer Wisconsin dataset from sklearn)

Adding a dataset

```
from sklearn.datasets import load_breast_cancer

class Dataset(BaseDataset):
    name = "Breast Cancer"

    def get_data(self):
        X, y = load_breast_cancer(return_X_y=True)
        return dict(X=X, y=y)
```

Adding metrics

```
class Objective(BaseObjective):
    name = "Classification"
    url = "https://github.com/#ORG/#BENCHMARK"
    requirements = ['scikit-learn']

    def set_data(self, X, y):
        self.X, self.y = X, y

    def evaluate_result(self, predict):
        return dict(
            accuracy_train=accuracy_score(self.y, predict(self.X)),
        )

    def get_objective(self):
        return dict(
            X_train=self.X,
            y_train=self.y,
        )
```

Adding metrics

```
class Objective(BaseObjective):
    name = "Classification"
    url = "https://github.com/#ORG/#BENCHMARK"
    requirements = ['scikit-learn']

    def set_data(self, X, y):
        self.X, self.y = X, y
        self.cv = KFold(
            n_splits=5, shuffle=True, random_state=self.get_seed()
        )
        self.cv_metadata = {}

    def evaluate_result(self, predict):
        return dict(
            accuracy_train=accuracy_score(self.y_train, predict(self.X_train)),
            accuracy_test=accuracy_score(self.y_test, predict(self.X_test)),
        )

    def get_objective(self):
        self.X_train, self.X_test, self.y_train, self.y_test = self.get_split(self.X, self.y)
        return dict(
            X_train=self.X_train,
            y_train=self.y_train,
        )
```

Adding metrics

Add a new metric (for example f1 score)

Adding metrics

```
from sklearn.metrics import f1_score

class Objective(BaseObjective):
    name = "Classification benchmark"
    url = "https://github.com/#ORG/#BENCHMARK"
    requirements = ['scikit-learn']

    def evaluate_result(self, predict):
        return dict(
            accuracy_train=accuracy_score(self.y_train, predict(self.X_train)),
            accuracy_test=accuracy_score(self.y_test, predict(self.X_test)),
            f1_test=f1_score(self.y_test, predict(self.X_test))
        )
```

Adding solvers

```
class Solver(BaseSolver):
    name = 'SVM'
    parameters = {
        'kernel': ['linear', 'poly', 'sigmoid'],
    }
    sampling_strategy = "run_once"

    def set_objective(self, X_train, y_train):
        self.X_train, self.y_train = X_train, y_train
        self.clf = SVC(kernel=self.kernel)

    def run(self, _):
        self.clf.fit(self.X_train, self.y_train)

    def get_result(self):
        return dict(predict=self.clf.predict)
```

Adding solvers

```
class Solver(BaseSolver):
    name = 'torch-logreg'
    parameters = {'lr': [0.01], 'batch_size': [64]}
    requirements = ['pytorch:pytorch']

    stopping_criterion = SufficientProgressCriterion(
        strategy="callback", eps=1e-3, patience=3,
        key_to_monitor="accuracy_test", minimize=False
    )

    def get_next(self, stop_val):
        return stop_val + 50

    def run(self, callback):
        optim = torch.optim.SGD(self.clf.parameters(), lr=self.lr)
        loss_fn = nn.CrossEntropyLoss()
        max_epochs = 100
        for _ in range(max_epochs):
            for X_batch, y_batch in self.dataloader:
                optim.zero_grad()
                y_pred = self.clf(X_batch)
                loss = loss_fn(y_pred, y_batch)
                loss.backward()
                optim.step()
                if not callback():
                    return
```

Adding solvers

Add a new solver (for example Logistic Regression from sklearn) or improve an existing one, by adding new parameters (optimiser, regularisation, batch size, etc.)

Adding plots

```
class Plot(BasePlot):
    name = "Accuracy"
    type = "bar_chart" # Or "scatter", "boxplot", "table"
    options = {"dataset": ..., "objective": ..., "split": ["train", "test"]}

    def plot(self, df, dataset, objective, split):
        df = df.query("dataset_name == @dataset and objective_name == @objective")
        plots = []
        for solver_name, df_filtered in df.groupby('solver_name'):
            df_filtered = df_filtered.select_dtypes(include=['number'])
            stop_val = df_filtered['stop_val'].max()
            this_df = df_filtered[df_filtered['stop_val'] == stop_val]
            plots.append({
                "y": this_df[f"accuracy_{split}"].tolist(),
                "text": "",
                "label": solver_name,
                "color": self.get_style(solver_name)["color"]
            })
        return plots

    def get_metadata(self, df, dataset, solver, objective, objective_column):
        return {
            "title": f"{objective}\nData: {dataset}\nSolver: {solver}",
            "ylabel": objective_column,
        }
```

Adding plots

More info on plots:

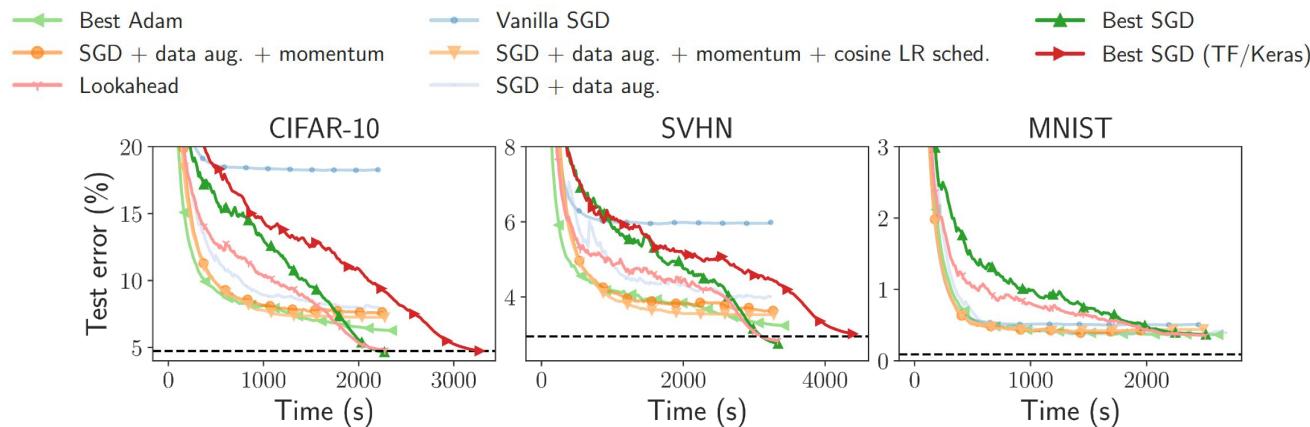
https://benchopt.github.io/stable/user_guide/add_custom_plot.html

Benchopt makes your life easy

- ▶ Build on previous benchmarks
- ▶ Use solvers in Python, R, Julia, binaries...
- ▶ Monitor any metric you want altogether (test/train loss, ...)
- ▶ Add parameters to solvers and use config files
- ▶ Share and publish HTML results
- ▶ Run all benchmarks in parallel, on HPC clusters...
- ▶ Cache runs and results
- ▶ and much more!

Example: Optimisation for ResNet on image classification

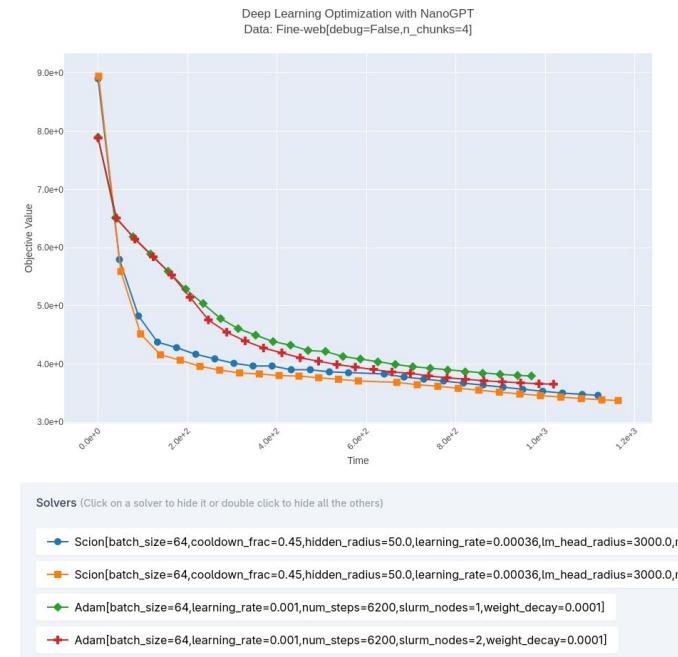
- Image classification with resnet18
- Various optimization strategies
- Compare pytorch and tensorflow
- Publish reproducible SOTA for baselines



https://github.com/benchopt/benchmark_resnet_classif/

Example: Large scale-optimisation for Deep learning

- Use modern large-scale datasets and models
- Classical optimization tricks
- Distributed training and mixed precision



https://github.com/benchopt/benchmark_nanogpt

Mini Project: Creating a benchmark

Some ideas:

1. MNIST classification

- Create the dataset (using MNIST from sklearn)
- Create deep learning solvers

2. Inverse Problems (deepinv)

- Create the datasets for a specific inverse problem
- Benchmark deepinv solvers

3. Optimization (Lasso)

- Create the datasets for regression or classification problems
- Benchmark different optimisers

You can also come up with your own ideas/benchmark, and we will also help you!